

CLOUD COMPUTING AND BIG DATA ANALYTICS

SAHANA S
(192571033)

1. Cloud Auto-Scaling and Load Balancing for Seasonal Traffic

Introduction

E-commerce websites often experience unpredictable traffic during events such as Black Friday, Diwali sales, Amazon Great Indian Festival, or Flipkart Big Billion Days. If the website has only a fixed number of servers, sudden traffic increases can overload the servers and cause slow response or downtime. Cloud auto-scaling and load balancing solve this problem.

Auto-Scaling

Auto-scaling automatically increases or decreases computing resources according to demand.

- During normal traffic, fewer virtual machines or containers are used.
- During a sudden traffic spike, additional instances are automatically created.
- When traffic decreases, unnecessary instances are removed.
- This ensures that resources are available when required.

For example:

Normal traffic → 3 servers

Sales traffic → 10 servers


After sale → 3 servers

This avoids both resource shortage and unnecessary infrastructure costs.

Load Balancing

A load balancer distributes incoming user requests across multiple servers.

For example:



Server 1 Server 2 Server 3

Instead of allowing all customers to connect to one server, requests are distributed among available servers.

Reduction of Downtime

Auto-scaling and load balancing reduce downtime by:

1. Preventing individual servers from becoming overloaded.
2. Automatically adding servers when demand increases.
3. Redirecting requests away from unhealthy servers.
4. Maintaining application availability during failures.
5. Supporting continuous service during traffic spikes.

Improvement in User Experience

Users experience:

- Faster page loading.
- Faster checkout.
- Fewer connection failures.
- Reduced waiting time.
- Better application responsiveness.

Role of Monitoring Tools

Monitoring tools continuously observe the cloud infrastructure.

They monitor:

- CPU utilization.
- Memory usage.
- Network traffic.
- Number of requests.
- Response time.
- Server health.
- Error rates.

For example, if CPU utilization increases above a predefined threshold such as 70%, the monitoring system can trigger an auto-scaling policy.

Real-World Example

Consider an e-commerce website during a Diwali sale.

Before the sale, the company may operate with 5 servers. When thousands of customers simultaneously search for products and place orders, traffic increases significantly.

The load balancer distributes requests among servers, while auto-scaling launches additional servers automatically.

After the sale ends, traffic decreases and the additional servers are removed.

Conclusion

Cloud auto-scaling and load balancing provide elasticity, high availability, cost efficiency, and better performance. Monitoring tools make the system intelligent by detecting changing workloads and triggering appropriate scaling actions.

2. Secure Cloud Storage for Backup Data

Introduction

Organizations generate huge amounts of backup data, including databases, application files, customer records, and business documents. Cloud storage provides a scalable and secure solution for storing this information.

Durability

Cloud providers store data using multiple physical storage systems.

Data can be replicated across:


- Multiple disks.
- Multiple servers.
- Multiple availability zones.
- Sometimes multiple geographical regions.

Therefore, failure of one storage device does not normally result in data loss.

Data Redundancy

Redundancy means maintaining multiple copies of data.

For example:



If one copy becomes unavailable, another copy can be used.

Encryption

Encryption protects backup data from unauthorized access.

Two important forms are:

1. Encryption at rest

Data is encrypted while stored in cloud storage.

2. Encryption in transit

Data is encrypted while moving between the organization and cloud storage.

Even if unauthorized users obtain the stored files, encrypted data is difficult to interpret without the required keys.

Access Control

Cloud platforms provide security mechanisms such as:

- Identity and Access Management (IAM).
- Role-based access control.
- Multi-factor authentication.
- Access policies.
- Least-privilege permissions.
- Audit logs.

For example, only the backup administrator may be allowed to delete backup files.

Cost Advantages

Traditional backup systems require organizations to purchase:

- Servers.
- Hard disks.
- Backup appliances.
- Data-center space.
- Cooling systems.
- Electricity.
- Maintenance.

Cloud storage reduces these capital expenses.

Cloud storage generally follows a pay-as-you-use model, allowing organizations to pay for the storage they actually consume.

Older or rarely accessed backups can also be moved to lower-cost archival storage.

Disaster Recovery Example

Suppose a company's primary data center is destroyed by a natural disaster.

The company can:

1. Recover backup data from cloud storage.
2. Restore databases and applications.
3. Start services in another region.
4. Continue business operations.

For example, an online banking company can maintain encrypted database backups in a separate cloud region. If its primary data center becomes unavailable, the backup can be restored at the disaster recovery site.

Conclusion

Cloud storage provides durability, redundancy, encryption, access control, scalability, and lower infrastructure costs, making it highly suitable for backup and disaster recovery.

3. CDN and Global Application Deployment

Introduction

When a mobile application is used by customers around the world, users located far away from the main application server may experience higher latency.

A Content Delivery Network (CDN) improves performance by distributing content through geographically distributed edge servers.

How CDN Works


Without CDN:

User India —————
User USA —————> Main Server

User Japan —————

All users communicate with the central server.

With CDN:



Users receive content from a nearby edge location.

Reduction in Latency

Latency is the time required for data to travel between the user and server.

If an Indian user requests content from a server located in the United States, the distance is large.

A CDN can serve cached content from an Indian or nearby edge server.

Therefore:

Shorter distance → Lower latency → Faster response

Edge Servers

Edge servers are geographically distributed servers located closer to users.

They can cache:

- Images.
- Videos.
- JavaScript files.
- CSS.
- Application data.
- Streaming content.

This reduces the need to repeatedly contact the origin server.

Cloud Global Availability

Cloud providers support global applications through:

- Multiple geographical regions.
- Availability zones.
- Global load balancing.
- CDN services.
- Replicated databases.
- Distributed storage.
- Disaster recovery mechanisms.

This allows applications to remain available even if one region experiences problems.

Real-Time Streaming Example

Consider a global live sports-streaming application.

During a cricket match, millions of users from India, the UK, Australia, and the USA may access the stream simultaneously.

The CDN can deliver video segments from nearby edge locations.

This results in:

- Lower latency.
- Faster video startup.
- Reduced buffering.
- Better video quality.
- Reduced load on the origin server.

Gaming Example

For an online multiplayer game, edge infrastructure can reduce the distance between players and game services.

This can improve:

- Response time.
- Gameplay responsiveness.
- Matchmaking.
- Download speeds.

Conclusion

CDNs and edge servers bring content closer to users, reducing latency and improving application performance. Combined with cloud regions, load balancing, and replication, they enable globally available and highly responsive applications.

4. Cloud Compliance and Security for Financial Institutions

Introduction

Financial institutions handle highly sensitive information such as:

- Bank account details.
- Credit card information.
- Customer identities.
- Transaction records.
- Financial statements.

Therefore, cloud environments must provide strong security, compliance, auditing, and data protection.

Regulatory Compliance

Cloud providers offer services and certifications that help organizations meet regulatory requirements.

Depending on the country and application, organizations may need to consider standards and regulations such as:

- PCI DSS.
- ISO 27001.
- SOC standards.
- GDPR where applicable.
- Local banking and data-protection regulations.

However, using a compliant cloud provider does not automatically make the customer's application compliant. The organization must configure and operate its systems correctly.

Logging

Cloud logging records important activities such as:

- User logins.
- Database access.
- File access.
- Configuration changes.
- API calls.
- Failed authentication attempts.

For example:

```
User → Login
      → Access Account
      → Transfer ₹50,000
      → Transaction Recorded
```

These events can be stored for auditing and investigation.

Monitoring

Monitoring continuously detects unusual activities.

It can identify:

- Repeated failed logins.
- Unusual transaction patterns.
- Unauthorized access.
- High network traffic.
- Suspicious API calls.

Security alerts can then be generated.

Encryption

Encryption protects sensitive information.

Data at rest: Financial records stored in databases are encrypted.

Data in transit: Transaction information travelling between systems is encrypted.

Encryption keys should be properly protected and access to them should be restricted.

Identity and Access Control

Financial organizations should implement:

- Multi-factor authentication.
- Role-based access.
- Least-privilege permissions.
- Strong password policies.
- Privileged-access controls.

For example, a normal employee should not have permission to modify transaction databases.

Data Sovereignty

Data sovereignty means that data may need to remain within a particular country's or jurisdiction's boundaries.

An organization can address this by:

1. Selecting cloud regions located in the required country.
2. Restricting storage locations using policies.
3. Controlling cross-region replication.
4. Reviewing provider contractual and compliance commitments.
5. Auditing where data is stored and processed.

Example: Online Banking Transaction

Suppose a customer transfers ₹1,00,000 through an online banking application.

The cloud system can:

1. Authenticate the customer using MFA.
2. Encrypt transaction information.
3. Record the transaction in audit logs.
4. Monitor for suspicious activity.

5. Apply access-control policies.
6. Store required financial data in approved geographical regions.

Conclusion

Cloud providers support financial organizations through encryption, identity management, logging, monitoring, compliance controls, auditing, and regional data controls. Proper configuration and governance are essential to protect financial transactions.

5. Containers, Virtual Machines and Kubernetes

Introduction

Containerization packages an application together with its dependencies so that it can run consistently across different environments.

Containers are widely used for modern microservices-based applications.

Containers vs Virtual Machines

Virtual Machines

A virtual machine includes:

Application
Libraries
Guest OS
Virtual Hardware
Hypervisor

Physical Hardware

Each VM generally requires its own guest operating system.

Containers

Containers share the host operating system kernel.

Application
Libraries
Container

Container Runtime

Host OS

Hardware

Comparison

Feature	Virtual Machine	Container
Operating system	Separate guest OS	Shares host kernel
Startup	Relatively slower	Very fast
Resource usage	Higher	Lower
Size	Usually larger	Usually smaller
Isolation	Strong	Process-level isolation
Portability	Good	Very high
Density	Fewer workloads per host	More workloads per host

Containers therefore allow organizations to run more application instances using fewer resources.

Kubernetes Orchestration

When an organization has hundreds of containers, manually managing them becomes difficult.

Kubernetes automates container management.

It can perform:

1. Deployment

Kubernetes automatically deploys containers across available nodes.

2. Scaling

If application traffic increases, Kubernetes can increase the number of running application instances.

Example:

Normal traffic → 3 Pods

High traffic → 10 Pods

3. Self-Healing

If a container fails, Kubernetes can automatically restart or replace it.

4. Load Balancing

Kubernetes can distribute requests among application instances.

5. Rolling Updates

New application versions can be deployed gradually without taking the entire application offline.

Portability Across Cloud Environments

Containers package the application and its dependencies into a standardized unit.


Therefore, the same container image can potentially run on:

- AWS.
- Microsoft Azure.
- Google Cloud.
- Private cloud.
- On-premises servers.

This reduces vendor lock-in and makes application migration easier.

Microservices Example

Consider an e-commerce application divided into microservices:



Each service can run in its own container.

For example:

- User Service → manages customer accounts.
- Product Service → manages products.
- Cart Service → manages shopping carts.
- Payment Service → processes payments.
- Order Service → manages orders.

Kubernetes manages these containers and automatically scales individual services according to demand.

Advantages

Containerization with Kubernetes provides:

- Faster deployment.
- Efficient resource utilization.
- Easy scaling.
- Self-healing.
- High availability.
- Portability.
- Consistent development and production environments.
- Easier microservices management.

Conclusion

Containers are lightweight compared with traditional virtual machines because they share the host OS kernel. Kubernetes adds automated deployment, scaling, load balancing, monitoring integration, and self-healing, making containers suitable for modern cloud-native applications.